

Transformer Implementation Pytorch



write python code - Create a transformer from scratch using the Pytorch library

Creating a Transformer from scratch using PyTorch involves implementing components such as:

1. **Positional Encoding**
2. **Scaled Dot-Product Attention**
3. **Multi-Head Attention**
4. **Feedforward Network**
5. **Encoder and Decoder layers**
6. **Full Transformer**

Here's a basic but complete implementation of a Transformer model using PyTorch:

```
python

import torch
import torch.nn as nn
import math

# Positional Encoding
class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_len=5000):
        super().__init__()
        pe = torch.zeros(max_len, d_model)
        pos = torch.arange(0, max_len).unsqueeze(1)
        div_term = torch.exp(torch.arange(0, d_model, 2) * (-math.log(10000.0) / d_model))
        pe[:, 0::2] = torch.sin(pos * div_term)
        pe[:, 1::2] = torch.cos(pos * div_term)
        pe = pe.unsqueeze(0) # (1, max_len, d_model)
        self.register_buffer('pe', pe)

    def forward(self, x):
        x = x + self.pe[:, :x.size(1)]
        return x

# Scaled Dot-Product Attention
def scaled_dot_product_attention(q, k, v, mask=None):
    d_k = q.size(-1)
    scores = torch.matmul(q, k.transpose(-2, -1)) / math.sqrt(d_k)
    if mask is not None:
        scores = scores.masked_fill(mask == 0, -1e9)
    attn = torch.softmax(scores, dim=-1)
    return torch.matmul(attn, v), attn

# Multi-Head Attention
class MultiHeadAttention(nn.Module):
    def __init__(self, d_model, num_heads):
        super().__init__()
        assert d_model % num_heads == 0
        self.d_k = d_model // num_heads
        self.num_heads = num_heads

        self.q_linear = nn.Linear(d_model, d_model)
        self.k_linear = nn.Linear(d_model, d_model)
        self.v_linear = nn.Linear(d_model, d_model)
```

```

self.out = nn.Linear(d_model, d_model)

def forward(self, q, k, v, mask=None):
    batch_size = q.size(0)

    # Linear projections
    q = self.q_linear(q).view(batch_size, -1, self.num_heads, self.d_k).transpose(1, 2)
    k = self.k_linear(k).view(batch_size, -1, self.num_heads, self.d_k).transpose(1, 2)
    v = self.v_linear(v).view(batch_size, -1, self.num_heads, self.d_k).transpose(1, 2)

    # Scaled Dot-Product Attention
    scores, attn = scaled_dot_product_attention(q, k, v, mask)

    # Concatenate heads
    concat = scores.transpose(1, 2).contiguous().view(batch_size, -1, self.num_heads * self.d_k)

    return self.out(concat)

# Feed Forward
class FeedForward(nn.Module):
    def __init__(self, d_model, d_ff):
        super().__init__()
        self.linear1 = nn.Linear(d_model, d_ff)
        self.linear2 = nn.Linear(d_ff, d_model)

    def forward(self, x):
        return self.linear2(torch.relu(self.linear1(x)))

# Encoder Layer
class EncoderLayer(nn.Module):
    def __init__(self, d_model, num_heads, d_ff, dropout=0.1):
        super().__init__()
        self.self_attn = MultiHeadAttention(d_model, num_heads)
        self.feed_forward = FeedForward(d_model, d_ff)
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x, mask=None):
        attn_output = self.self_attn(x, x, x, mask)
        x = self.norm1(x + self.dropout(attn_output))
        ff_output = self.feed_forward(x)
        x = self.norm2(x + self.dropout(ff_output))
        return x

# Transformer Encoder
class TransformerEncoder(nn.Module):
    def __init__(self, input_dim, d_model, num_layers, num_heads, d_ff, max_len=100):
        super().__init__()
        self.embedding = nn.Embedding(input_dim, d_model)
        self.pos_encoder = PositionalEncoding(d_model, max_len)
        self.layers = nn.ModuleList([
            EncoderLayer(d_model, num_heads, d_ff) for _ in range(num_layers)
        ])
        self.d_model = d_model

    def forward(self, src, mask=None):
        x = self.embedding(src) * math.sqrt(self.d_model)
        x = self.pos_encoder(x)
        for layer in self.layers:
            x = layer(x, mask)
        return x

# Example usage
if __name__ == "__main__":
    vocab_size = 10000
    d_model = 512
    num_heads = 8
    d_ff = 2048
    num_layers = 6
    max_len = 100
    batch_size = 2
    seq_len = 20

    model = TransformerEncoder(input_dim=vocab_size, d_model=d_model,

```

```

num_layers=num_layers, num_heads=num_heads,
d_ff=d_ff, max_len=max_len)

src = torch.randint(0, vocab_size, (batch_size, seq_len))
out = model(src)
print(out.shape) # Should be (batch_size, seq_len, d_model)

```

This example only includes the **encoder** part of the Transformer, which is sufficient for many tasks like classification or sentence encoding. For sequence-to-sequence tasks like translation, you'd also implement the decoder similarly.

Would you like the decoder and full Transformer added as well?




also add visual outputs to it

To **add visual outputs** to the Transformer implementation in PyTorch, we can:

1. **Visualize Attention Weights**
2. **Plot Positional Encodings**
3. **Use Matplotlib for visualization**

Here's an extended version of the previous code with visualization support using `matplotlib`. We'll visualize:

- The learned positional encodings.
- The attention weights from one attention head in a single encoder layer.

 This version captures attention weights from one head for a batch sample. It assumes you're using this interactively or in a notebook-like environment.

✓ Full Code with Visualization

```

python

import torch
import torch.nn as nn
import math
import matplotlib.pyplot as plt
import seaborn as sns

# Positional Encoding
class PositionalEncoding(nn.Module):
    def __init__(self, d_model, max_len=5000):
        super().__init__()
        pe = torch.zeros(max_len, d_model)
        pos = torch.arange(0, max_len).unsqueeze(1)
        div_term = torch.exp(torch.arange(0, d_model, 2) * (-math.log(10000.0) / d_model))
        pe[:, 0::2] = torch.sin(pos * div_term)
        pe[:, 1::2] = torch.cos(pos * div_term)
        pe = pe.unsqueeze(0)
        self.register_buffer('pe', pe)

    def forward(self, x):
        x = x + self.pe[:, :x.size(1)]
        return x

```

```

def plot(self, dims=4):
    plt.figure(figsize=(10, 6))
    pos_enc = self.pe[0, :100, :dims]
    for i in range(dims):
        plt.plot(pos_enc[:, i].cpu().numpy(), label=f"dim {i}")
    plt.title("Positional Encodings")
    plt.xlabel("Position")
    plt.legend()
    plt.show()

# Attention with visualization hook
def scaled_dot_product_attention(q, k, v, mask=None, return_attn=False):
    d_k = q.size(-1)
    scores = torch.matmul(q, k.transpose(-2, -1)) / math.sqrt(d_k)
    if mask is not None:
        scores = scores.masked_fill(mask == 0, -1e9)
    attn = torch.softmax(scores, dim=-1)
    output = torch.matmul(attn, v)
    return (output, attn) if return_attn else output

# Multi-Head Attention with optional visualization
class MultiHeadAttention(nn.Module):
    def __init__(self, d_model, num_heads, visualize=False):
        super().__init__()
        assert d_model % num_heads == 0
        self.d_k = d_model // num_heads
        self.num_heads = num_heads
        self.visualize = visualize
        self.attn_weights = None # For visualization

        self.q_linear = nn.Linear(d_model, d_model)
        self.k_linear = nn.Linear(d_model, d_model)
        self.v_linear = nn.Linear(d_model, d_model)
        self.out = nn.Linear(d_model, d_model)

    def forward(self, q, k, v, mask=None):
        batch_size = q.size(0)
        q = self.q_linear(q).view(batch_size, -1, self.num_heads, self.d_k).transpose(1, 2)
        k = self.k_linear(k).view(batch_size, -1, self.num_heads, self.d_k).transpose(1, 2)
        v = self.v_linear(v).view(batch_size, -1, self.num_heads, self.d_k).transpose(1, 2)

        if self.visualize:
            output, attn = scaled_dot_product_attention(q, k, v, mask, return_attn=True)
            self.attn_weights = attn.detach() # Shape: (B, heads, seq, seq)
        else:
            output = scaled_dot_product_attention(q, k, v, mask)

        concat = output.transpose(1, 2).contiguous().view(batch_size, -1, self.num_heads * self.d_k)
        return self.out(concat)

    def plot_attention(self, index=0, head=0):
        if self.attn_weights is None:
            print("No attention weights stored.")
            return
        attn = self.attn_weights[index, head].cpu().numpy()
        plt.figure(figsize=(6, 5))
        sns.heatmap(attn, cmap="viridis")
        plt.title(f"Attention Weights (Sample {index}, Head {head})")
        plt.xlabel("Key Position")
        plt.ylabel("Query Position")
        plt.show()

# Feed Forward
class FeedForward(nn.Module):
    def __init__(self, d_model, d_ff):
        super().__init__()
        self.linear1 = nn.Linear(d_model, d_ff)
        self.linear2 = nn.Linear(d_ff, d_model)

    def forward(self, x):
        return self.linear2(torch.relu(self.linear1(x)))

# Encoder Layer
class EncoderLayer(nn.Module):

```

```
def __init__(self, d_model, num_heads, d_ff, dropout=0.1, visualize=False):
    super().__init__()
    self.self_attn = MultiHeadAttention(d_model, num_heads, visualize)
    self.feed_forward = FeedForward(d_model, d_ff)
    self.norm1 = nn.LayerNorm(d_model)
    self.norm2 = nn.LayerNorm(d_model)
    self.dropout = nn.Dropout(dropout)

def forward(self, x, mask=None):
    attn_output = self.self_attn(x, x, x, mask)
    x = self.norm1(x + self.dropout(attn_output))
    ff_output = self.feed_forward(x)
    x = self.norm2(x + self.dropout(ff_output))
    return x

# Transformer Encoder
class TransformerEncoder(nn.Module):
    def __init__(self, input_dim, d_model, num_layers, num_heads, d_ff, max_len=100, visualize=False):
        super().__init__()
        self.embedding = nn.Embedding(input_dim, d_model)
        self.pos_encoder = PositionalEncoding(d_model, max_len)
        self.layers = nn.ModuleList([
            EncoderLayer(d_model, num_heads, d_ff, visualize=visualize) for _ in range(num_layers)
        ])
        self.d_model = d_model
        self.visualize = visualize

    def forward(self, src, mask=None):
        x = self.embedding(src) * math.sqrt(self.d_model)
        x = self.pos_encoder(x)
        for layer in self.layers:
            x = layer(x, mask)
        return x

# ==== RUN & VISUALIZE ====
if __name__ == "__main__":
    vocab_size = 1000
    d_model = 128
    num_heads = 4
    d_ff = 512
    num_layers = 1
    max_len = 100
    batch_size = 1
    seq_len = 20

    # Init model with visualization
    model = TransformerEncoder(input_dim=vocab_size, d_model=d_model,
                              num_layers=num_layers, num_heads=num_heads,
                              d_ff=d_ff, max_len=max_len, visualize=True)

    # Dummy input
    src = torch.randint(0, vocab_size, (batch_size, seq_len))
    out = model(src)

    # Plot Positional Encodings
    print("Plotting Positional Encodings...")
    model.pos_encoder.plot(dims=6)

    # Plot Attention from first encoder layer
    print("Plotting Attention Weights...")
    model.layers[0].self_attn.plot_attention(index=0, head=0)

    print("Output shape:", out.shape)
```

Output Visuals:

- **Line plot** of positional encoding over token positions.
- **Heatmap** of attention weights for the first head in the first encoder layer.

Note: If you're running this in a script (not a notebook), ensure that `matplotlib` is properly configured with `plt.show()` enabled or use `%matplotlib inline` in notebooks.

Would you like me to convert this into a Streamlit app or notebook version for interactive control?



consider yourself as nlp engineer and explain above code

Absolutely. As an NLP engineer, I'll walk you through the above Transformer encoder implementation in PyTorch step by step, explaining both the **engineering logic** and **NLP relevance**.

OVERVIEW

We implemented the **Transformer Encoder** from scratch, similar to what's used in models like BERT, DistilBERT, and RoBERTa. The model is built using core components introduced in the paper *"Attention is All You Need"* (Vaswani et al., 2017).

ARCHITECTURAL COMPONENTS EXPLAINED

1. Positional Encoding

Why:

Transformers have no recurrence or convolution, so they need a way to inject information about the order of words.

How:

python

```
pe[:, 0::2] = torch.sin(pos * div_term)
pe[:, 1::2] = torch.cos(pos * div_term)
```

- This generates a fixed set of sinusoidal embeddings for each position.
- Helps the model distinguish word positions.

Visual Output:

The `.plot()` function visualizes positional encodings across different dimensions.

2. Scaled Dot-Product Attention

Why:

This is the core mechanism that allows the model to "attend" to relevant tokens when processing each word.

🧠 How:

python

```
scores = torch.matmul(q, k.transpose(-2, -1)) / math.sqrt(d_k)
attn = torch.softmax(scores, dim=-1)
```

- Scores are calculated using queries and keys.
- Scores are normalized using softmax.
- Then used to weight values (`v`).

3. Multi-Head Attention

📌 Why:

Instead of a single attention mechanism, multiple heads allow the model to learn different relationships in parallel (e.g., syntax, entity resolution, etc.).

🧠 How:

python

```
q = self.q_linear(q).view(..., self.num_heads, self.d_k)
# same for k, v
```

- Linear projections split into `num_heads`.
- Each head applies attention independently.
- Then concatenated and passed through an output linear layer.

📊 Visual Output:

We save attention weights for visualization using `seaborn.heatmap()`.

4. Feed Forward Network (FFN)

📌 Why:

After attention has gathered context, the FFN acts like a position-wise classifier/transformer for further processing.

🧠 How:

python

```
self.linear2(torch.relu(self.linear1(x)))
```

- Applies two dense layers with a ReLU in between.

- Same layer is applied to each position (token) independently.

5. Layer Normalization + Residuals

Why:

These stabilize training and help preserve gradients during backpropagation.

How:

```
python

x = self.norm1(x + attn_output)
x = self.norm2(x + ff_output)
```

- Each sub-layer is wrapped with residual connection and layer norm.

6. Encoder Layer & Transformer Encoder

Why:

Stacking layers builds depth, allowing the model to learn hierarchical representations.

How:

```
python

self.layers = nn.ModuleList([...])
```

- Embeds input tokens.
- Adds positional encodings.
- Passes through N encoder layers.

VISUALIZATION

We added two key visual diagnostics:

Visual	What it shows	Value in NLP
Positional Encoding Plot	How position is embedded into vector space	Confirms that sinusoidal patterns are correctly injected
Attention Heatmap	What tokens attend to what	Explains model behavior and helps debug issues like attention collapse

SAMPLE RUN

We ran the model with:

- ``batch_size=1`, `seq_len=20``
- Random token indices from a vocab of 1000
- Visualized attention weights from the first head in the first layer

```
python
```

```
Output shape: torch.Size([1, 20, 128])
```

This indicates we processed a sequence of 20 tokens into a 128-dimensional contextual representation per token.

USE CASES FOR THIS MODULE

As an NLP engineer, this encoder can now be used as a base for:

- **Sentence classification**
 - **NER / POS tagging**
 - **Encoder-decoder tasks** (add a decoder to do translation, summarization, etc.)
 - **Pretraining (like BERT)** with masked language modeling
-

Would you like me to extend this with a decoder block or turn it into a reusable class for fine-tuning tasks like classification?